

CS164 Final Projects

Viet Hoang Tran Duong

Fall 2019

Problem Statement: We receive tons of emails every day, many of them are spam emails. How can we filter the spam email, using the concepts learned from this course?

Approach: We try two approaches and compare the accuracy of them: Logistic Regression and Support Vector Machine (SVM), both using the cvxpy package.

Constraints: Logistic Regression and SVM try to optimize the defined loss function, with the decision variable θ , which is a coefficient of a linear function. By nature, θ does not have constraints as it runs freely in R^n .

I will add layers of constraints by adding the l_1 and l_∞ Regularization. To convert the norms into linear optimization, I introduce slack variables as constraints, which will be discussed below.

Data: The data is from Kaggle SMS Spam Collection Dataset, which contains 5572 messages, tagged according to being ham (legitimate) or spam.

The data is subset into train set (80%) and test set (20%) randomly by sklearn train_test_split.

Data Processing:

Step 1: Remove punctuations and stopwords: Because the data is in texts (natural language), we need to preprocess it to fit into the optimization models. First, I remove the punctuation and stop words ("stopwords" from nltk package).

Step 2: Vectorize the text: To vectorize the texts, first, I use the term frequency-inverse document frequency method. However, it yields more dimensions than my data (9000+ dimensions > 5572 data points). Hence, I use a better approach: python spacy package "en_core_web_lg." This package will map the SMS content into the 300-dimension vector space.

As we have little data, 300 dimensions are still too large, which might lead to overfitting the train set. Hence, I apply Principal Component Analysis (sklearn PCA) to reduce the dimensionality while retaining as much variance in the data as possible. I selected the top 30 dimensions that capture most of the variance in the data. Our data becomes (5572, 30): 5572 data points in 30 dimensions.

Step 3: Transform the type decode for 2 algorithms: Because Logistic Regression works best with {0,1} while SVM works better with {-1,1}. I map the {ham, spam} type to these values.

Outline for the methods:

- Logistic Regression: 3 versions: the standard version, with l_1 regularization, and with l_∞ regularization.
- SVM: 3 versions: the standard version, with l_1 regularization, and with l_∞ regularization.

All methods are then evaluated by the train set and test set errors.

Regularization: Introducing the norm (l_1, l_∞) into the cost functions will punish the factor that causes overfit the train data. l_1 norm will push the coefficient to be towards 0: reducing the number of features in the model altogether. l_∞ will punish the most expensive coefficient factor (the biggest absolute value). Overall, regularization avoids overfitting, and subsequently, increases the accuracy.

Logistic Regression: Input: X (features matrix), Y (desired output), Output θ : coefficient
There are m data points, each with n features (in the data, it's 5572 data points and 30 features).

Feasible region: $\theta \in R^n$

Decision function: $\text{sigmoid}(x) = \frac{1}{1+e^{-\theta x}}$: the *sigmoid* function will return the value between 0 and 1: which is the probability of the email is spam. If the value is larger than 0.5, then return 1 (spam), else, return 0 (not spam).

We will minimize the cost the function of the log-likelihood function of the data:

$$L(\theta) = -\frac{1}{m} \sum_{i=1}^m (y_i \theta x_i - \ln(1 + e^{\theta x_i}))$$

We will prove that $L(\theta)$ is convex:

$$\begin{aligned} g(\theta) &= \ln(1 + e^{\theta x}) \\ \Rightarrow g'(\theta) &= \frac{(1+e^{\theta x})'}{1+e^{\theta x}} = \frac{x e^{\theta x}}{1+e^{\theta x}} = \frac{x}{1+e^{-\theta x}} \\ \Rightarrow g''(\theta) &= \frac{x'(1+e^{-\theta x}) - (1+e^{-\theta x})' x}{(1+e^{-\theta x})^2} = \frac{0(1+e^{-\theta x}) - (-x)e^{-\theta x} x}{(1+e^{-\theta x})^2} = \frac{x^2 e^{-\theta x}}{(1+e^{-\theta x})^2} \geq 0 \quad \forall \theta, x \in D \end{aligned}$$

(because $x^2 \geq 0$, $e^{-\theta x} \geq 0$, $(1 + e^{-\theta x})^2 \geq 0$)

As the second-order derivative of $g(x)$ is always nonnegative, hence, $g(x)$ is convex

We have:

- $y_i \theta x_i$ is a linear function: affine \Rightarrow concave
- $g(\theta) = \ln(1 + e^{\theta x})$: convex $\forall \theta, x \Rightarrow -\ln(1 + e^{\theta x_i})$: concave $\forall \theta, x_i$
- Sum of concave functions is a concave function

$\Rightarrow \sum_{i=1}^m (y_i \theta x_i - \ln(1 + e^{\theta x_i}))$: concave function. Multiply with $-\frac{1}{m} < 0$ will turn this into convex:

$$\Rightarrow L(\theta) = -\frac{1}{m} \sum_{i=1}^m (y_i \theta x_i - \ln(1 + e^{\theta x_i})) \text{ : convex function.} \quad (*1)$$

Hence, this problem is a convex optimization problem.

To avoid overfitting, I will perform regularization with a l_1 and l_∞ norm: penalize the loss function if the norms become too big. The regularization term is $\lambda|\theta|$, with λ is a set parameter.

For l_1 : $\Rightarrow L^*(\theta) = -\frac{1}{m} \sum_{i=1}^m (y_i \theta x_i - \ln(1 + e^{\theta x_i})) + \lambda|\theta|$

We introduce a slack variable: $s, s \in R^{n \times 1}$ such that: $|\theta| < s \Leftrightarrow -s \leq \theta \leq s$

The loss function becomes: $L^*(\theta, s) = -\frac{1}{m} \sum_{i=1}^m (y_i \theta x_i - \ln(1 + e^{\theta x_i})) + \lambda 1^T s$

As $\lambda 1^T s$ is a linear function: affine \Rightarrow convex, $L(\theta)$ is a convex function (*1), hence, the sum of these 2 functions: $L^*(\theta, s) = -\frac{1}{m} \sum_{i=1}^m (y_i \theta x_i - \ln(1 + e^{\theta x_i})) + \lambda 1^T s$: convex. Therefore, even with the regularization term, this is still a convex optimization problem. (*2)

For l_∞ : $\Rightarrow L^*(\theta) = -\frac{1}{m} \sum_{i=1}^m (y_i \theta x_i - \ln(1 + e^{\theta x_i})) + \lambda |\theta|_\infty$

We introduce a slack variable: $s, s \in R$ such that: $\max(|\theta|) < s \Leftrightarrow |\theta_i| \leq s \forall i \Leftrightarrow 1 * -s \leq \theta \leq 1 * s$
The loss function becomes: $L^*(\theta, s) = -\frac{1}{m} \sum_{i=1}^m (y_i \theta x_i - \ln(1 + e^{\theta x_i})) + \lambda s$

As λs is a linear function: affine \Rightarrow convex, $L(\theta)$ is a convex function (*1), hence, the sum of these 2 functions: $L^*(\theta, s) = -\frac{1}{m} \sum_{i=1}^m (y_i \theta x_i - \ln(1 + e^{\theta x_i})) + \lambda s$: convex. Therefore, even with the regularization term, this is still a convex optimization problem. (*3)

From (*1), (*2), (*3): All the loss function (with/without regularization) are all convex, and we are trying to minimize them. Hence, it is guaranteed to find the optimal solutions.

Without the regularization term, here are the results:

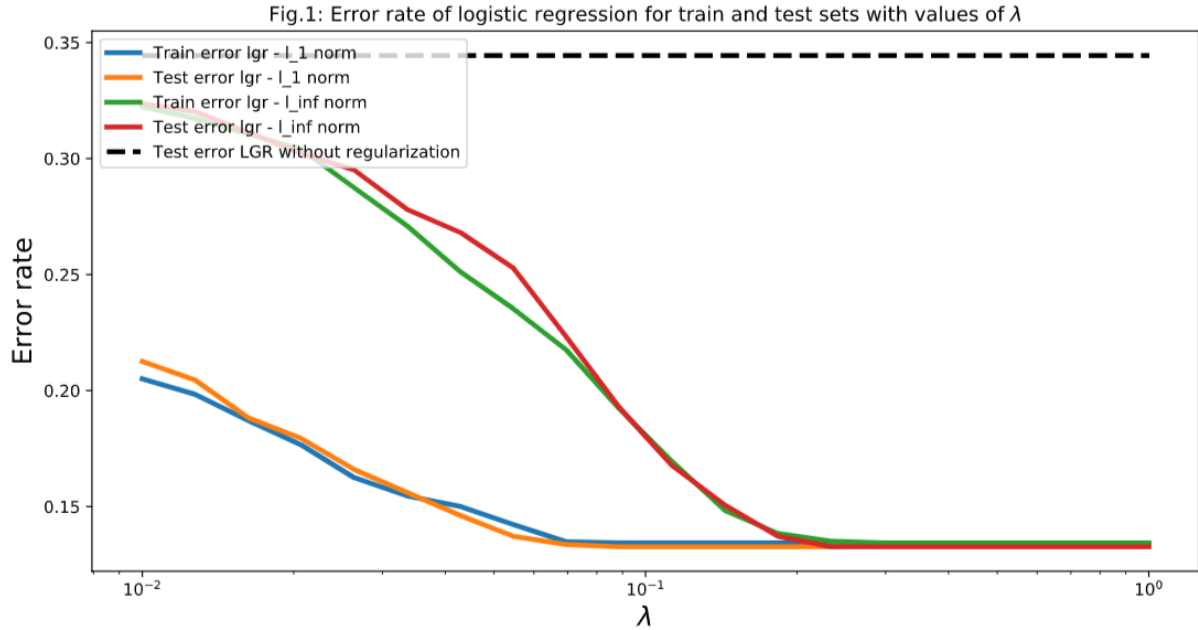
Results for logistic regression without regularization:

Status: optimal

Train error: 0.34305586717522996

Test error: 0.34439461883408073

With the regularization term: I vary the value of λ with the 10 exponents of 20 uniformly distributed values in $[-2, 0]$: $\lambda \in [10^{-2}, 10^0]$



SVM: Input: X (features matrix), Y (desired output); **Output** θ : coefficient

There are m data points, each with n features (in the data, it's 5572 data points and 30 features).

Feasible region: $\theta \in R^n$

Decision function: $y = \text{sign}(\theta x - v)$: if $\theta x - v > 0$: $y = 1$ (spam); else and $y = -1$ (not spam)

Objective function: minimize loss function:

$$L(\theta, v) = \frac{1}{m} \sum_{i=1}^m (1 - y_i(\theta x_i - v))$$

We have, $L(\theta, v)$ is a linear function, hence, convex.

(*4)

Hence, this is a convex optimization problem.

To avoid overfitting, I will perform regularization with a l_1 and l_∞ norm: penalize the loss function if the norms become too big. The regularization term is $\lambda|\theta|$, with λ is a set parameter.

For l_1 : $\Rightarrow L^*(\theta, v) = \frac{1}{m} \sum_{i=1}^m (1 - y_i(\theta x_i - v)) + \lambda|\theta|_1$

We introduce a slack variable: $s, s \in R^{n \times 1}$ such that: $|\theta|_1 < s \Leftrightarrow -s \leq \theta \leq s$

The loss function becomes: $L^*(\theta, v, s) = \frac{1}{m} \sum_{i=1}^m (1 - y_i(\theta x_i - v)) + \lambda 1^T s$

As $\lambda 1^T s$ is a linear function: affine \Rightarrow convex, $L(\theta)$ is a convex function (*), hence, the sum of these 2 functions: $L^*(\theta, v, s) = \frac{1}{m} \sum_{i=1}^m (1 - y_i(\theta x_i - v)) + \lambda 1^T s$: convex. Therefore, even with the regularization term, this is still a convex optimization problem.

(*5)

For l_∞ : $\Rightarrow L^*(\theta, v) = \frac{1}{m} \sum_{i=1}^m (1 - y_i(\theta x_i - v)) + \lambda|\theta|_\infty$

We introduce a slack variable: $s, s \in R$ such that: $\max(|\theta|) < s \Leftrightarrow |\theta_i| \leq s \forall i \Leftrightarrow 1 * -s \leq \theta \leq 1 * s$

The loss function becomes: $L^*(\theta, s) = \frac{1}{m} \sum_{i=1}^m (1 - y_i(\theta x_i - v)) + \lambda s$

As λs is a linear function: affine \Rightarrow convex, $L(\theta)$ is a convex function (*), hence, the sum of these 2 functions: $L^*(\theta, s) = \frac{1}{m} \sum_{i=1}^m (1 - y_i(\theta x_i - v)) + \lambda s$: convex. Therefore, even with the regularization term, this is still a convex optimization problem.

(*6)

From (*4), (*5), (*6): All the loss function (with/without regularization) are all convex, and we are trying to minimize them, it is guaranteed to find the optimal solutions.

Without the regularization term, here are the results:


Results for SVM without regularization:

Status: optimal

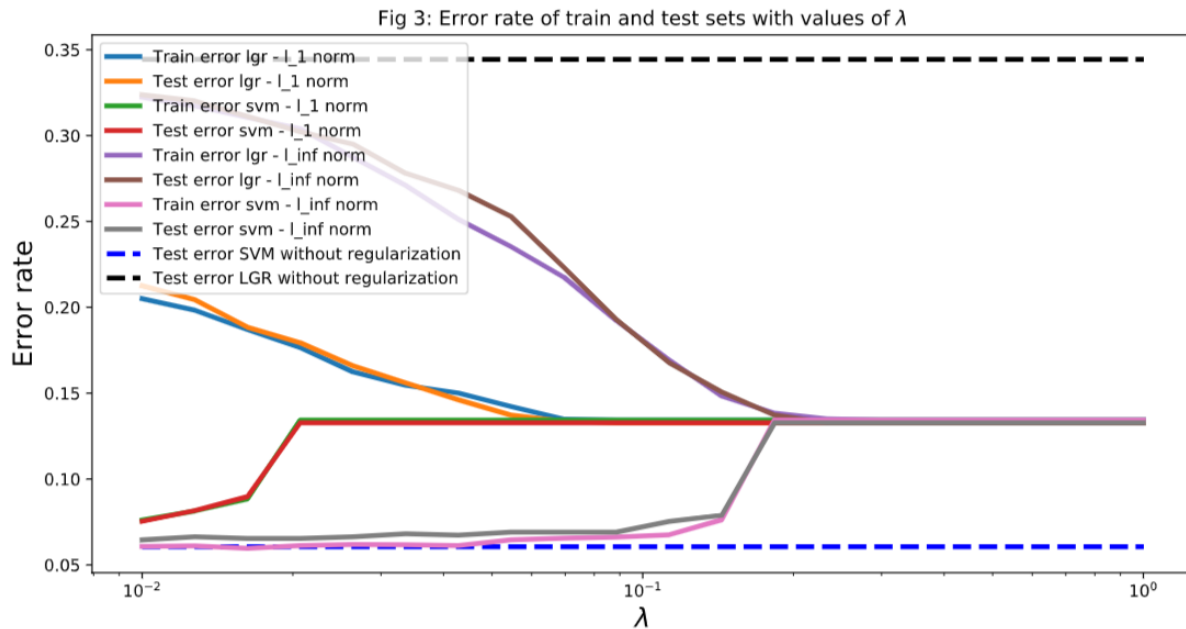
Train error: 0.060578864707202154

Test error: 0.060578864707202154

With the regularization term, I vary the value of λ with the 10 exponent of 20 uniformly distributed values in $[-2,0]$: $\lambda \in [10^{-2}, 10^0]$



Total results:



Conclusion:

Overall, these classic methods can be solved by convex optimization, as the nature of the loss functions is convex. Furthermore, we can add regularization through slack variables, which still maintain the convex property (while adding linear constraints).

SVM outperformed the Logistic Regression, especially when $\lambda = 10^{-2}$. For logistics regression, increasing λ increases the accuracy, while increasing λ makes the performance of SVM worse.

HCS applications:

#context: I tailored my writing to fit the context of a mathematics and code report, which are concise, on point, adding with multiple illustrations (figures) to compare between methods. This paper has demonstrated proper use of #context, where I understand the context of this assignment and build content upon.

#biasidentification + #biasmitigation: I apply the concept of bias into the data science: the bias-variance tradeoff: I apply regularization to reduce the variance of the estimator by simplifying it (discourage learning a more complex model by shrinking the coefficient estimates towards 0), which will avoid the risk of overfitting. The bias in Empirical Analysis relates to bias-variance reduction for a better model, which is essentially achieved by reducing overfitting using regularization. This paper is an excellent application of #biasidentification (acknowledging the existence of high variance), and fixing the overfitting by using #biasmitigation,

#connotation: The connotation (tone, style) of the language within this document is clear, concise, with correct terminology and methods application. All these components contribute to a concise report that suitable for the targeted audience: those interested in simple methods (logistics, SVM) to solve abstract natural language problems.

LOs applications:

#convexprogramming: For this project, the objective functions are all convex, even with or without the regularization factor. I also implement all the code in python, with combined use of numpy, cvxpy, pandas, and other packages. Finally, I automate the system to solve multiply convex optimizations problems to evaluate the parameters' performance.

#convexity: I analytically show the loss function of the logistic regression (the log-likelihood function) is convex by using the second-order derivative. I also analytically showed all functions to be convex, turning the norm factors into linear programming using slack variables. These proofs are demonstrations of the proper application of #convexity.

References:

- Dataman. (2019, October 24). Dimension Reduction Techniques with Python. Retrieved from <https://towardsdatascience.com/dimension-reduction-techniques-with-python-f36ca7009e5c>.
- Loh, P.-S. (n.d.). Convexity - math.cmu.edu. Retrieved from http://www.math.cmu.edu/~ploh/public_html/docs/math/mop2013/convexity-soln.pdf.
- MLMath.io. (2019, February 16). Math behind Support Vector Machine(SVM). Retrieved from <https://medium.com/@ankitnitjsr13/math-behind-support-vector-machine-svm-5e7376d0ee4d>.
- Sanjeevi, M., & Mady. (2017, September 26). Chapter 2.0 : Logistic Regression with Math. Retrieved from <https://medium.com/deep-math-machine-learning-ai/chapter-2-0-logistic-regression-with-math-e9cbb3ec6077>.
- UCI Machine Learning. (2016, December 2). SMS Spam Collection Dataset. Retrieved from <https://www.kaggle.com/uciml/sms-spam-collection-dataset/data>.

Link to code:

- https://colab.research.google.com/drive/1nusjx9nli2fO_ri0_2f7Ybruzw5ZF5Ai
- Addition: 4 layers Neural Network built from scratch train on the same dataset:
https://colab.research.google.com/drive/1Q0BvysgFPHhuK91Vb5llqAegNsZ_m0zL

```
1 # import packages: numpy, pandas, nltk, spacy, and sklearn
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import pandas as pd
5 import spacy
6 from sklearn.decomposition import PCA
7 import time
8 import nltk
9 nltk.download('stopwords')
10 from nltk.corpus import stopwords
11 from sklearn.model_selection import train_test_split
12 import cvxpy as cvx
13 import random
14 import copy
15 import string
16 %matplotlib inline
17 %config InlineBackend.figure_format = 'svg'
18
19 # set seed for consistency
20 np.random.seed(2019)
```

```
↳ [nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
1 # package installation
2
3 #!spacy download en_core_web_lg
```

```
1 # dimension setting function
2 # if this code fell to run (can't recognize package en_core_web_lg), then run the li
3 # !spacy download en_core_web_lg
4 # then restart the kernel.
5
6 nlp = spacy.load('en_core_web_lg')
```

▼ Read the data

```

1 # spam dataset: https://www.kaggle.com/uciml/sms-spam-collection-dataset/data
2 data = pd.read_csv("https://docs.google.com/spreadsheets/d/e/2PACX-1vRzkLSK-20SdsUII
3
4 # only takes the 1st 2 columns (type and content)
5 data = data[["v1", "v2"]]
6 # rename the columns
7 data.columns = ["Type", "Content"]
8
9 # check null values
10 for column in data.columns:
11     if sum(data[column].isnull()) != 0:
12         print("Exist Null value: ", column)
13
14 print("Number of data points:", len(data))
15 print("")
16 data.head()

```

➞ Number of data points: 5572

	Type	Content
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

```

1 # quick description on the data
2 data.groupby("Type").describe()

```

➞

Content

	count	unique	top	freq
Type				
ham	4825	4516	Sorry, I'll call later	30
spam	747	647	Please call our customer service representativ...	4

▼ Data Preprocessing

```

1 # remove punctuations and stopwords
2
3 def txt_preprocess(txt):
4     txt = txt.translate(str.maketrans('', '', string.punctuation))
5     txt = [word for word in txt.split() if word.lower() not in stopwords.words('engli:
6     return " ".join(txt)
7
8 # make a copy of the data so that we won't change the original dataframe
9 process_text = copy.deepcopy(data["Content"])
10
11 # preprocess the data
12 process_text = process_text.apply(txt_preprocess)

```

```

1 # map the data into 300 dimensions spacy vector space
2 train_nlp = [nlp(text) for text in process_text]
3 train_vecs = [text.vector for text in train_nlp]
4
5 # feature matrix
6 X = np.array(train_vecs)

```

```

1 # use pca to remove the dimension of the data: 300 -> 30 dimensions: keep those with
2 pca = PCA(n_components=30)
3 X = pca.fit_transform(X)

```

```

1 # transform the type data: logistic regression {0,1}; SVM {-1,1}
2 Y_lgr = np.array([data["Type"].map({"ham": 0, 'spam':1})]).T
3 Y_svm = np.array([data["Type"].map({"ham": -1, 'spam':1})]).T

```

```

1 # split the dataset into train and test
2 X_train, X_test, Y_train_lgr, Y_test_lgr = train_test_split(X,
3                                                             Y_lgr, test_size=0.2, random
4
5 # this code is to have consistency between train and test (for both approaches to h
6 # logistics: {0,1} * 2 --> {0,2} - 1 --> {-1,1}: the values suitable for SVM
7 Y_train_svm = Y_train_lgr*2-1
8 Y_test_svm = Y_test_lgr*2-1

```

```

1 # check dimension
2 m = X_train.shape[0]
3 n = X_train.shape[1]
4 ones = np.ones((n,1))
5
6 print(m, "data points")
7 print(n, "features")

```

```

↳ 4457 data points
   30 features

```

```

1

```

logistic with L₁ norm

```

1 # error check
2 def error(scores, labels):
3     return np.mean((np.sign(scores + 1)//2) != labels)

```

```

1 # define variables
2 coeff = cvx.Variable((n,1))
3 slack_1 = cvx.Variable((n,1), nonneg = True)
4

```

```

5 # Parameters
6 lambda = cvx.Parameter(nonneg=True)
7
8 # cost function
9 cost = cvx.sum(cvx.multiply(Y_train_lgr, X_train @ coeff) - cvx.logistic(X_train @ (
10
11 # objective function: -1/m*cost + regularization
12 obj = cvx.Minimize(-cost/m + lambda * (ones.T@slack_1))
13
14 # constraints
15 con = [coeff >= -slack_1, coeff <= slack_1]
16
17 # define problem
18 problem = cvx.Problem(obj, con)

```

```

1 # Try out with no regularization: lambda = 0
2 lambda.value = 0
3 problem.solve(solver = cvx.ECOS)
4 print("Results for logistic regression without regularization:")
5 print("Status:", problem.status)
6
7 # calculate train/test error
8 train_error_lgr = error((X_train @ coeff).value, Y_train_lgr)
9 test_error_lgr = error((X_test @ coeff).value, Y_test_lgr)
10
11 print("Train error:", train_error_lgr)
12 print("Test error:", test_error_lgr)

```

```

↳ Results for logistic regression without regularization:
   Status: optimal
   Train error: 0.34305586717522996
   Test error: 0.34439461883408073

```

```


1 # number of values for lambda:
2 trials = 20
3
4 # to store the error rate
5 train_error_lgr_1 = []

```

```
6 test_error_lgr_1 = []
7
8 # space out the values by log of 10: to 10^-2 to 10^0
9 lambda_vals = np.logspace(-2, 0, trials)
10
11 # try out with different lambda values
12 for i in range(trials):
13     lambd.value = lambda_vals[i]
14     problem.solve(solver = cvx.ECOS)
15
16 # store the error rate
17 train_error_lgr_1 += [error((X_train @ coeff).value, Y_train_lgr)]
18 test_error_lgr_1 += [error((X_test @ coeff).value, Y_test_lgr)]

1 # plot the error rates
2 plt.figure(figsize=(12,6))
3 plt.plot(lambda_vals, train_error_lgr_1, label="Train error lgr - l_1 norm", linewidth=2)
4 plt.plot(lambda_vals, test_error_lgr_1, label="Test error lgr - l_1 norm", linewidth=2)
5 plt.xscale("log")
6 plt.legend(loc="upper left")
7 plt.xlabel(r"$\lambda$", fontsize=16)
8 plt.ylabel("Error rate", fontsize=16)
9 plt.show()
```





1

logistic with linf norm

```
1 # variables
2 coeff = cvx.Variable((n,1))
3 slack_inf = cvx.Variable((1,1))
```



```

4
5 # define parameters
6 lambda = cvx.Parameter(nonneg=True)
7
8 # cost function
9 cost = cvx.sum(cvx.multiply(Y_train_lgr, X_train @ coeff) - cvx.logistic(X_train @ c
10
11 # objective function
12 obj = cvx.Minimize(-cost/m + lambda * slack_inf)
13
14 # constraint for slack variable
15 con = [coeff >= -slack_inf@ones, coeff <= slack_inf@ones]
16
17 # define problem
18 problem = cvx.Problem(obj, con)

```

```

1 # number of values for lambda:
2 trials = 20
3
4 # to store the error rate
5 train_error_lgr_inf = []
6 test_error_lgr_inf = []
7
8 # space out the values by log of 10: to 10^-2 to 10^0
9 lambda_vals = np.logspace(-2, 0, trials)
10
11 # try out with different lambda values
12 for i in range(trials):
13     lambda.value = lambda_vals[i]
14     problem.solve(solver = cvx.ECOS)
15
16     # store the error rate
17     train_error_lgr_inf += [error((X_train @ coeff).value, Y_train_lgr)]
18     test_error_lgr_inf += [error((X_test @ coeff).value, Y_test_lgr)]


```

```

1 # plot the error rates
2 plt.figure(figsize=(12,6))
3 plt.plot(lambda_vals, train_error_lgr_inf, label="Train error lgr - l_inf norm", lir

```


```
4 plt.plot(lambda_vals, test_error_lgr_int, label="Test error lgr - l_inf norm", linev
5 plt.xscale("log")
6 plt.legend(loc="upper left")
7 plt.xlabel(r"$\lambda$", fontsize=16)
8 plt.ylabel("Error rate", fontsize=16)
9 plt.show()
```



Logistic summary

```
1 # plot the combines error rates
2 plt.figure(figsize=(12,6))
3 plt.plot(lambda_vals, train_error_lgr_1, label="Train error lgr - l_1 norm", linewidth=2)
4 plt.plot(lambda_vals, test_error_lgr_1, label="Test error lgr - l_1 norm", linewidth=2)
5 plt.plot(lambda_vals, train_error_lgr_inf, label="Train error lgr - l_inf norm", linewidth=2)
6 plt.plot(lambda_vals, test_error_lgr_inf, label="Test error lgr - l_inf norm", linewidth=2)
7 plt.hlines(test_error_lgr, 0.01, 1, linestyle = "--", label="Test error LGR without regularization")
8 plt.xscale('log')
9 plt.legend(loc='upper left')
10 plt.xlabel(r"$\lambda$", fontsize=16)
11 plt.ylabel("Error rate", fontsize=16)
12 plt.title("Fig.1: Error rate of logistic regression for train and test sets with varying lambda")
13 plt.show()
```



Fig.1: Error rate of logistic regression for train and test sets with values of λ 

1

SVM with L1 norm

```
1 # define variables
2 coeff = cvx.Variable((n,1))
```

```

3 v = cvx.Variable()
4 slack_1 = cvx.Variable((n,1), nonneg = True)
5
6 # cost functions
7 cost = cvx.sum(cvx.pos(1 - cvx.multiply(Y_train_svm, X_train*coeff - v)))
8
9 # parameters
10 lambd = cvx.Parameter(nonneg=True)
11
12 # objective function
13 obj = cvx.Minimize(cost/m + lambd * (ones.T@slack_1))
14
15 # constraint
16 con = [coeff >= -slack_1, coeff <= slack_1]
17
18 # define problem
19 problem = cvx.Problem(obj, con)

```

```

1 # try out with out regularization
2 lambd.value = 0
3 problem.solve(solver = cvx.ECOS)
4 print("Results for SVM without regularization:")
5 print("Status:", problem.status)
6
7 train_error_svm = np.mean(Y_train_svm != np.sign(X_train.dot(coeff.value) - v.value))
8 test_error_svm = np.mean(Y_train_svm != np.sign(X_train.dot(coeff.value) - v.value))
9
10 print("Train error:", train_error_svm)
11 print("Test error:", test_error_svm)

```

```

↳ Results for SVM without regularization:
   Status: optimal
   Train error: 0.060578864707202154
   Test error: 0.060578864707202154

```

```

1 # number of values for lambd:
2 trials = 20
3
4 # to store the error rate

```

```

7 # store the error rate
5 train_error_svm_1 = []
6 test_error_svm_1 = []
7
8 # space out the values by log of 10: to 10^-2 to 10^0
9 lambda_vals = np.logspace(-2, 0, trials)
10
11 # try out with different lambda values
12 for i in range(trials):
13     lambda.value = lambda_vals[i]
14     problem.solve(solver = cvx.ECOS)
15
16     # store the error rate
17     train_error_svm_1 += [np.mean(Y_train_svm != np.sign(X_train.dot(coeff.value) -
18     test_error_svm_1 += [np.mean(Y_test_svm != np.sign(X_test.dot(coeff.value) - v.\


```

```

1 plt.figure(figsize=(12,6))
2 plt.plot(lambda_vals, train_error_svm_1, label="Train error svm - l_1 norm", linewidth=2)
3 plt.plot(lambda_vals, test_error_svm_1, label="Test error svm - l_1 norm", linewidth=2)
4 plt.xscale('log')
5 plt.legend(loc='upper left')
6 plt.xlabel(r"$\lambda$", fontsize=16)
7 plt.ylabel("Error rate", fontsize=16)
8 plt.show()

```





1

SVM linf

```
1 # define variable
2 coeff = cvx.Variable((n,1))
3 v = cvx.Variable()
```

```

4 slack_inf = cvx.variable((1,1))
5
6 # cost function
7 cost = cvx.sum(cvx.pos(1 - cvx.multiply(Y_train_svm, X_train*coeff - v)))
8
9 # parameters
10 lambd = cvx.Parameter(nonneg=True)
11
12 # objective function
13 obj = cvx.Minimize(cost/m + lambd * slack_inf)
14
15 # constraint
16 con = [coeff >= -slack_inf@ones, coeff <= slack_inf@ones]
17
18 # define problem
19 problem = cvx.Problem(obj, con)

```

```

1 # number of values for lambd:
2 trials = 20
3
4 # to store the error rate
5 train_error_svm_inf = []
6 test_error_svm_inf = []
7
8 # space out the values by log of 10: to 10^-2 to 10^0
9 lambda_vals = np.logspace(-2, 0, trials)
10
11 # try out with different lambd values
12 for i in range(trials):
13     lambd.value = lambda_vals[i]
14     problem.solve(solver = cvx.ECOS)
15
16     # store the error rate
17     train_error_svm_inf += [np.mean(Y_train_svm != np.sign(X_train.dot(coeff.value))
18     test_error_svm_inf += [np.mean(Y_test_svm != np.sign(X_test.dot(coeff.value)) - \

```

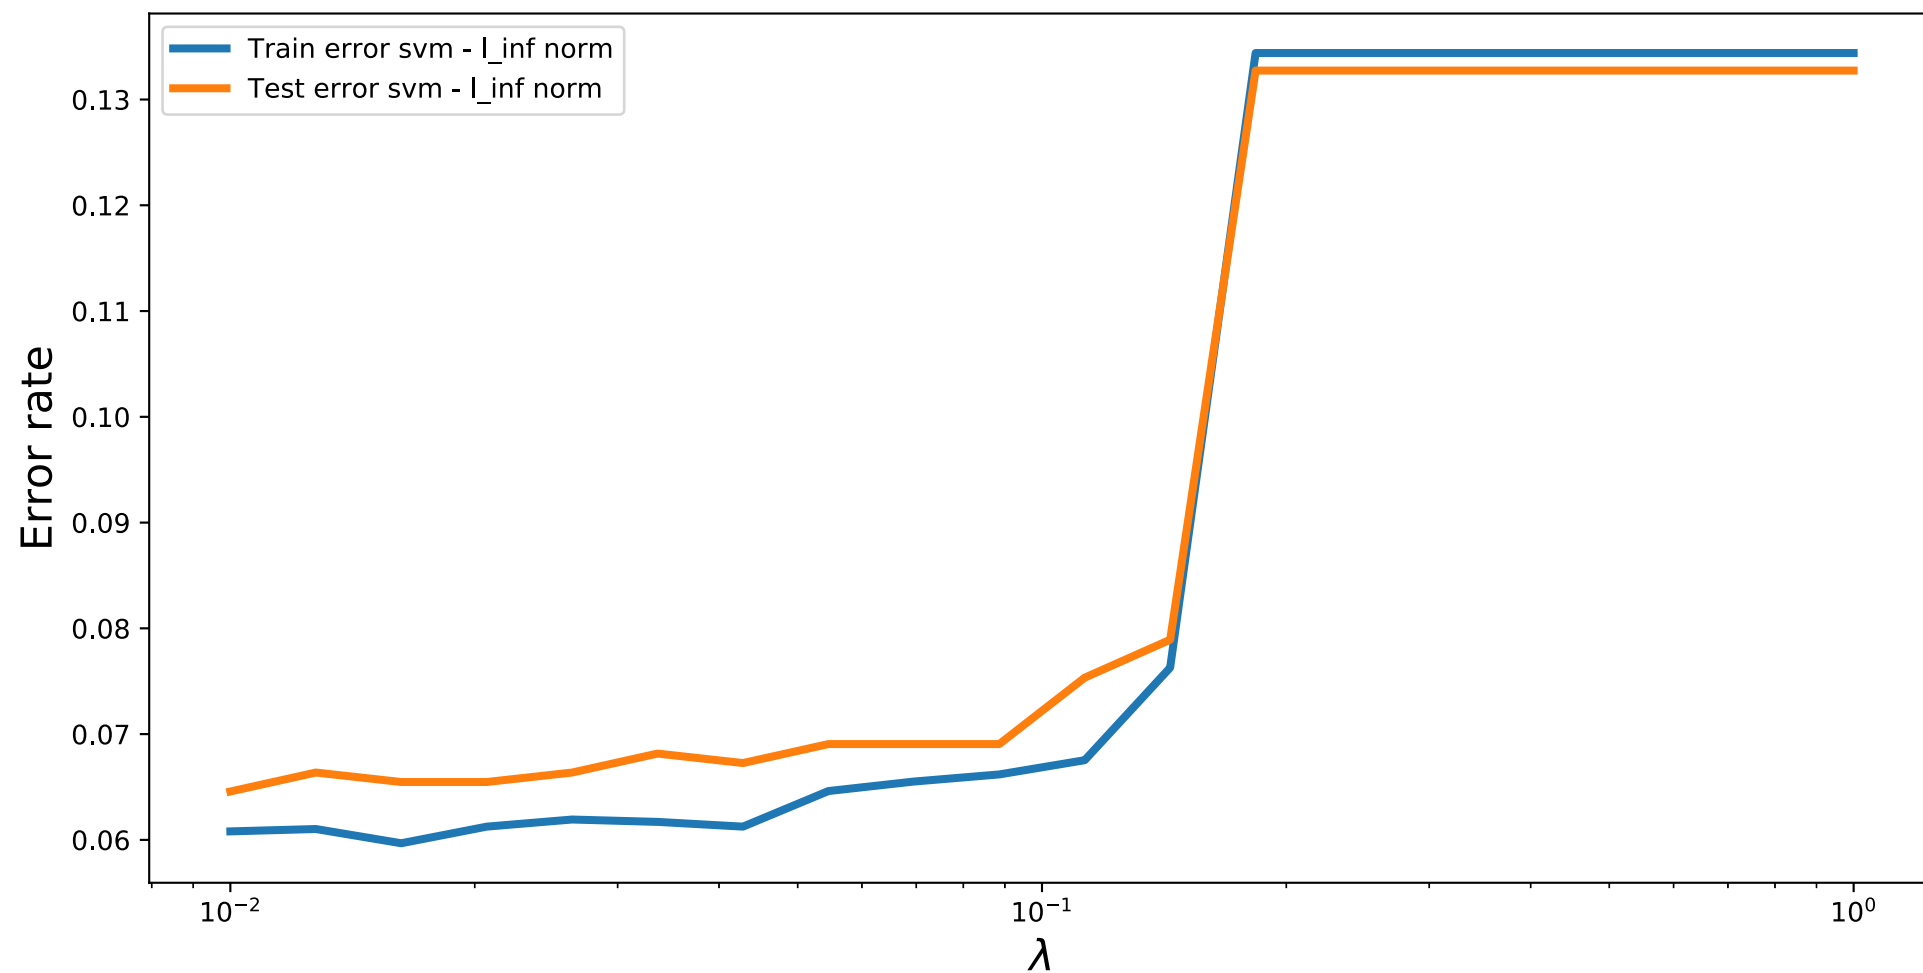
```

1 plt.figure(figsize=(12,6))
2 plt.plot(lambda_vals, train_error_svm_inf, label="Train error svm - l_inf norm", lir

```



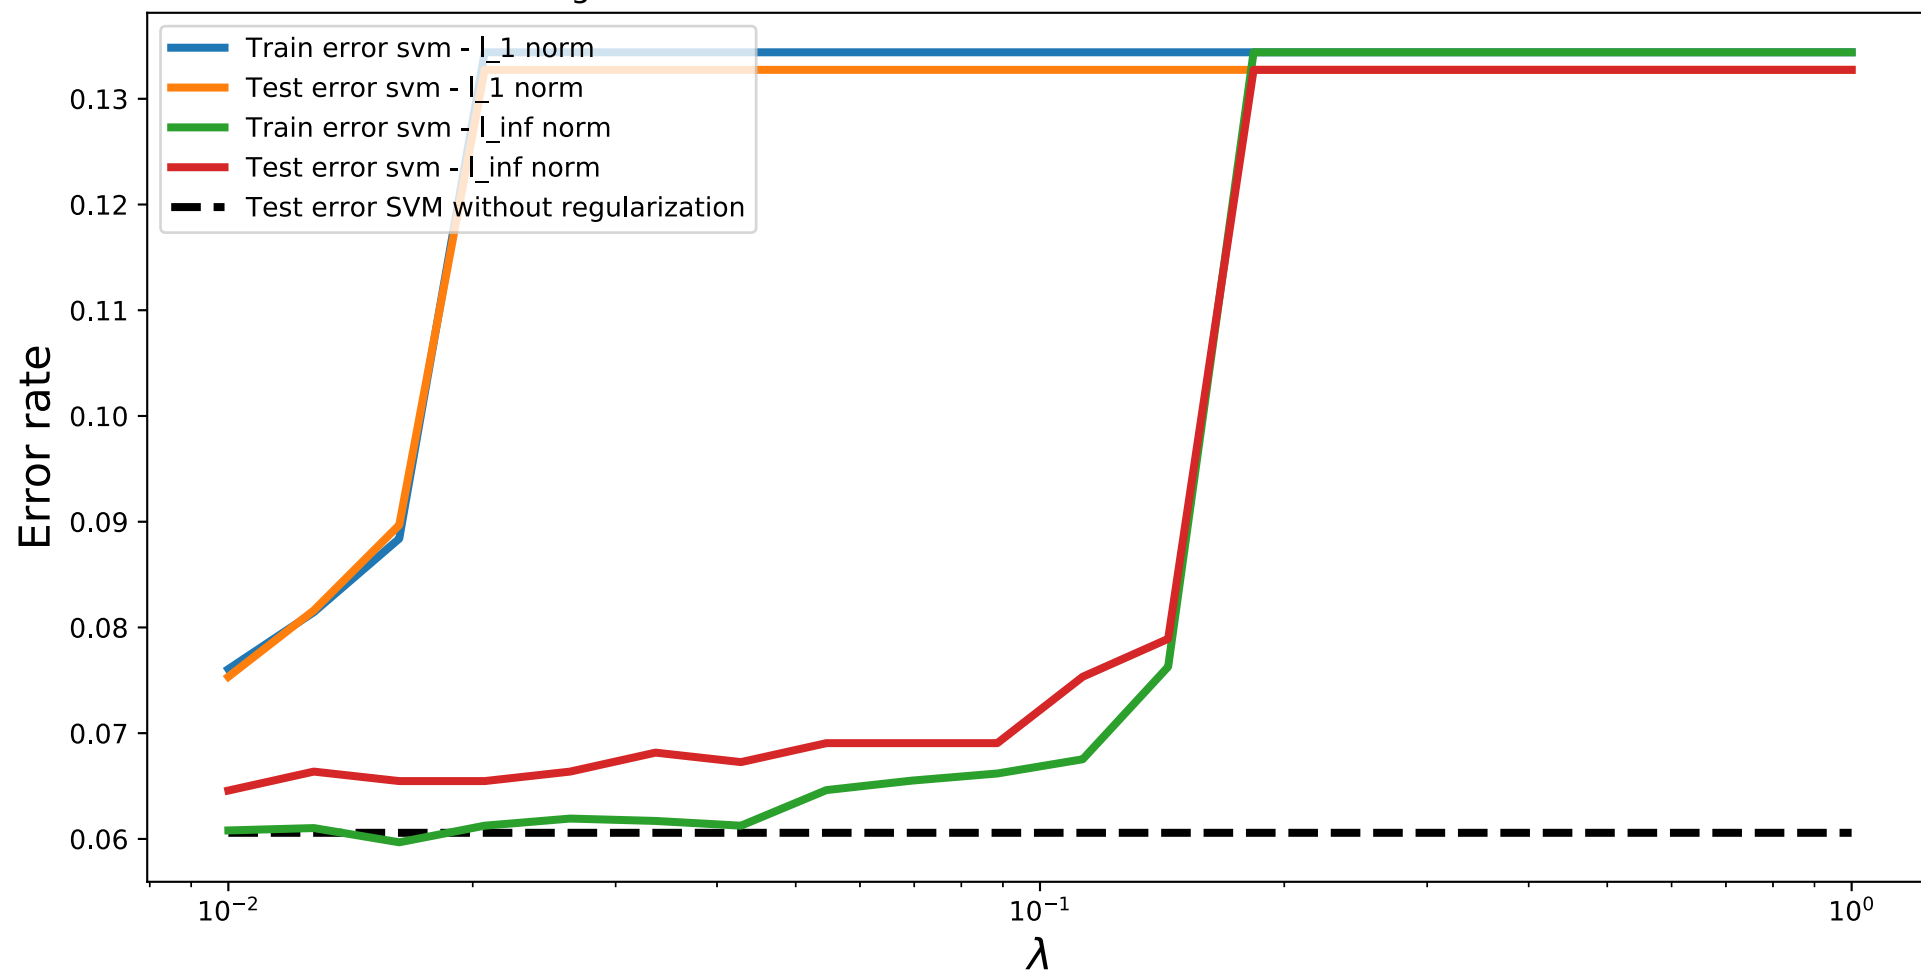
```
3 plt.plot(lambda_vals, test_error_svm_inf, label='Test error svm - l_inf norm', line  
4 plt.xscale('log')  
5 plt.legend(loc='upper left')  
6 plt.xlabel(r"$\lambda$", fontsize=16)  
7 plt.ylabel("Error rate", fontsize=16)  
8 plt.show()
```



SVM Summary

```
1 plt.figure(figsize=(12,6))
2 plt.plot(lambda_vals, train_error_svm_1, label="Train error svm - l_1 norm", linewidth=2)
3 plt.plot(lambda_vals, test_error_svm_1, label="Test error svm - l_1 norm", linewidth=2)
4 plt.plot(lambda_vals, train_error_svm_inf, label="Train error svm - l_inf norm", linewidth=2)
5 plt.plot(lambda_vals, test_error_svm_inf, label="Test error svm - l_inf norm", linewidth=2)
6 plt.hlines(test_error_svm, 0.01, 1, linestyle = "--", label="Test error SVM without l1")
7 plt.xscale('log')
8 plt.legend(loc='upper left')
9 plt.xlabel(r"$\lambda$", fontsize=16)
10 plt.ylabel("Error rate", fontsize=16)
11 plt.title("Fig 2: Error rate of SVM for train and test sets with values of $\lambda$")
12 plt.show()
```



Fig 2: Error rate of SVM for train and test sets with values of λ 

1

Total results

```
1 plt.figure(figsize=(12,6))
2 plt.plot(lambda_vals, train_error_lgr_1, label="Train error lgr - l1 norm", linewidth=2)
3 plt.plot(lambda_vals, test_error_lgr_1, label="Test error lgr - l1 norm", linewidth=2)
4 plt.plot(lambda_vals, train_error_lgr_2, label="Train error lgr - l2 norm", linewidth=2)
5 plt.plot(lambda_vals, test_error_lgr_2, label="Test error lgr - l2 norm", linewidth=2)
6 plt.plot(lambda_vals, test_error_svm_no_reg, label="Test error SVM without regularization", linewidth=2, linestyle='dashed')
7 plt.legend()
8 plt.show()
```

```
3 plt.plot(lambda_vals, test_error_lgr_1, label="Test error lgr - l_1 norm", linewidth=2)
4 plt.plot(lambda_vals, train_error_svm_1, label="Train error svm - l_1 norm", linewidth=2)
5 plt.plot(lambda_vals, test_error_svm_1, label="Test error svm - l_1 norm", linewidth=2)
6 plt.plot(lambda_vals, train_error_lgr_inf, label="Train error lgr - l_inf norm", linewidth=2)
7 plt.plot(lambda_vals, test_error_lgr_inf, label="Test error lgr - l_inf norm", linewidth=2)
8 plt.plot(lambda_vals, train_error_svm_inf, label="Train error svm - l_inf norm", linewidth=2)
9 plt.plot(lambda_vals, test_error_svm_inf, label="Test error svm - l_inf norm", linewidth=2)
10 plt.hlines(test_error_svm, 0.01, 1, linestyle = "--", label="Test error SVM without l1")
11 plt.hlines(test_error_lgr, 0.01, 1, linestyle = "--", label="Test error LGR without l1")
12 plt.xscale('log')
13 plt.legend(loc='upper left')
14 plt.xlabel(r"$\lambda$", fontsize=16)
15 plt.ylabel("Error rate", fontsize=16)
16 plt.title("Fig 3: Error rate of train and test sets with values of $\lambda$")
17 plt.show()
```



Fig 3: Error rate of train and test sets with values of λ 